Flight Delays Prediction

Optimization via Grid Search and Ensemble Models

Q. LAPPARRA, S. MARTIN BOUSSER, N.LE BRIS

ESILV

November 26, 2025

Presentation Outline

- **Objective:** Binary Classification (Predict `ARR_DEL15`).
  - Class 0: No Delay (< 15 min) | Class 1: Delay (≥ 15 min).
- **Data Source:** Kaggle `Jan_2019_ontime.csv` (reduced to 10% / 56,597 rows).
- **Why the "753 Variables" ?:**
  - **Core Variables ($\approx$ 10):** Date (Month/Day), Time (Blocks), Carrier, Origin, Dest, Distance.
  - **One-Hot Encoding impact:**
    - Airports ($\approx$ 353 origins + 353 dests) $\rightarrow$ 706 columns.
    - Carriers ($\approx$ 17) & Time Blocks ($\approx$ 19) $\rightarrow$ 36 columns.
  - **Result:** A matrix of **753 binary features**.

Airline Performance Analysis

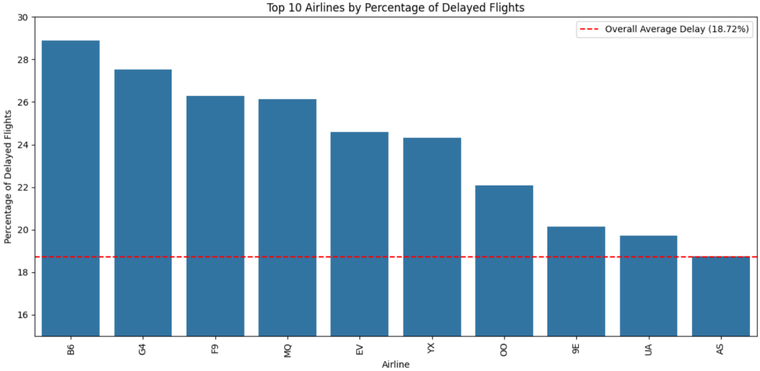

Figure: Top 10 Airlines by Percentage of Delayed Flights. The red line represents the overall average delay ($\approx 18.72\%$).

- Airlines like **B6** and **G4** show delay rates significantly above average.
- Carrier identity is a strong predictor feature.

Data Refinery: Feature Importance

Refining the 753 Variables:
Out of the massive feature space

generated by One-Hot Encoding, we identified which signals truly drive the model.

Key Insights:

- **Carrier Identity (`OP_CARRIER`):** The single most important predictor. Who flies the plane matters more than where it goes.
- **Logistics:** `DISTANCE` and `DEP_TIME` are critical structural factors.
- **Granularity:** Individual airports (e.g., SFO, HOU) appear in the top 10, but

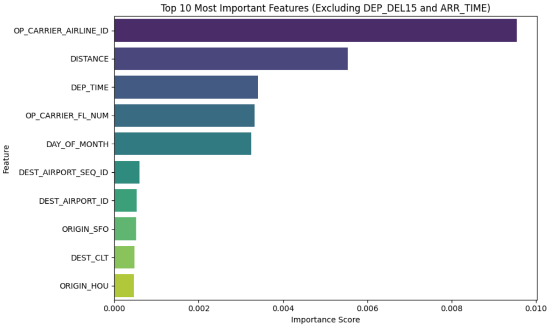

Figure: Top 10 Most Important Features (Excluding `DEP_DEL15`).

Baseline vs. Tuned Models

Initial Decision Tree:

- Accuracy: 87.42%
- High False Negatives (647).

After Grid Search:

- `max_depth`: 2
- Accuracy: $\approx$ 92.06%
- Specificity improved significantly.

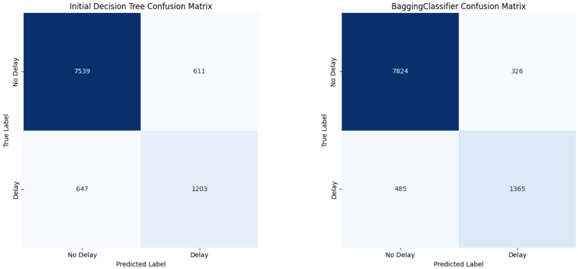

Figure: Comparison: Initial Decision Tree (Left) vs. Bagging/Tuned Model (Right).

Hyperparameter Tuning Results

We optimized the Decision Tree using `GridSearchCV` (CV=2).

Optimal Parameters:

- `criterion`: 'gini'
- `max_depth`: **2** (Preventing overfitting).
- `min_samples_split`: 2

Key Finding: The Grid Search revealed that a simpler model performs better than a complex one for this specific feature set.

Bagging Classifier: ROC Analysis

Model Performance:

- **Accuracy:** 91.89%
- **AUC Score:** 0.89

Interpretation: The ROC curve (Receiver Operating Characteristic) shows a strong ability to distinguish between classes. An AUC of 0.89 is considered excellent for this type of noisy real-world data.

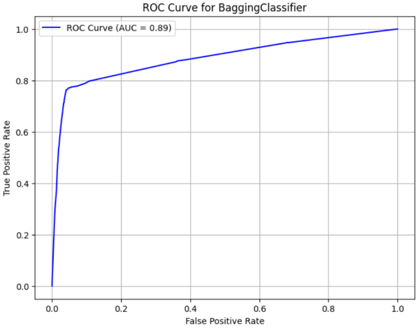

Figure: ROC Curve for Bagging Classifier.

Bagging Classifier: Confusion Matrix

Confusion Matrix Analysis:

- **True Negatives:** 7824 (Excellent prediction of on-time flights).
- **False Negatives:** 485 (Delays missed).

The model is conservative: it prefers predicting "No Delay" to avoid False Positives, which explains the high accuracy but slightly lower Recall on delays.

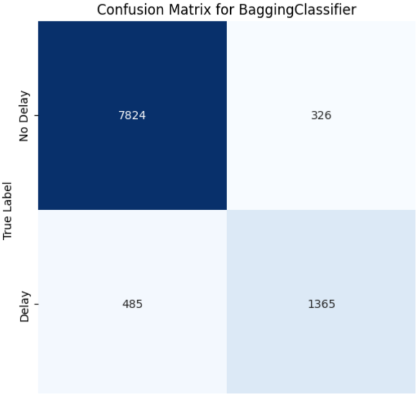

Figure: Detailed Confusion Matrix.

Key Takeaways Next Steps

Conclusion:

- Feature space exploded to 753 dims due to the implementation of Airports/Carriers.
- A simple model (`max_depth=2`) achieved the best balance ($\approx 92\%$).
- Bagging confirmed the stability of the model (AUC 0.89).

Next Steps:

- **Feature Engineering:** Incorporate weather data.
- **Advanced Models:** Test LightGBM.

References

Divyansh22. *Flight Delay Prediction*. Kaggle Dataset.

Thank you!